

Input Encoders for Multi-Channel Signal Transformers: An Empirical Audit of Subspace Geometry

Ossi Lehtinen*

Ocon Oy

Abstract

Transformers consuming multi-channel scalar signals must embed C simultaneous values into one d_{model} -dimensional vector per time step. We empirically audit eight input encoders—spanning a shared-scalar baseline, per-channel linear projections, an orthogonality regulariser, a nonlinear MLP stem, block-partitioned concatenation, channel-independent and channel-as-token architectures, and a projected positional encoding—on a synthetic benchmark designed to make channel identity informative and on ETTh1 as a real-data check. Most comparisons are run at 20 paired seeds and reported in negative log-likelihood (NLL) of the next-step quantile bin; validation top-1 accuracy is tracked alongside and tells a consistent secondary story. The headline is one of practical near-equivalence within a wide “top tier”: the standard per-channel linear projection (`nn.Linear(C, dmodel)`) matches every alternative we test up to small, statistically real, but practically modest differences. The only encoder that decisively loses is the shared-scalar baseline, which collapses for information-theoretic reasons we make explicit. Paired tests do resolve two small gaps: projecting the sinusoidal positional encoding through a learned linear layer edges the rest at small C , with a direct geometric probe showing the mechanism is positional-channel orthogonalisation; a nonlinear MLP stem edges them at the largest C we test, with the gap shrinking under more training data. The ordering is unchanged when the categorical head is swapped for an MSE regression head. The practical recommendation is to use `nn.Linear(C, dmodel)` by default and reach for something more elaborate only when the task at hand gives a real reason to do so. Code: <https://github.com/OssiLehtinen/FDM-encoding>.

1 Introduction

Transformers applied to multivariate numerical signals—multi-sensor telemetry, financial factor stacks, biomedical waveforms—face a design choice at the input layer: how to embed C simultaneous scalar channels into one d_{model} -dimensional vector per time step.

Prior art. The additive embedding recipe inherited from NLP (1) projects each input to a vector and sums. When the inputs are multiple channels rather than one token-plus-position pair, this comes in two forms: a *shared* scalar projection W with per-channel additive embeddings, or a *per-channel* projection W_k —the latter being mathematically a single `nn.Linear(C, dmodel)`. Many multivariate time-series transformers use a linear input stem (2, 3), though details vary. Alternative architectures treat each channel as a separate token (iTransformer (5)), run independent backbones per channel (PatchTST (4)), or combine both (Crossformer (6)). Orthogonality penalties on weight matrices have been used for feature decorrelation in deep networks (7).

*ossi@ocon.fi

Open questions. Despite the variety of approaches, several basic questions remain unanswered in the literature: (i) Compared to a bare `nn.Linear(C,  $d_{\text{model}}$ )` projection, what do we actually gain—in held-out loss, in geometric structure, or in convergence speed—from block partitioning, an orthogonality penalty, a nonlinear MLP stem, or a channel-independent / channel-as-token architecture? (ii) Does the task loss itself produce per-channel separation without explicit enforcement? (iii) Does the positional encoding interact with the channel encoding, and can this interaction be improved? We include a shared-scalar baseline (SUM) as a verified floor rather than an open question—the encoder algebraically collapses C channels to their pointwise sum, so its failure mode is information-theoretic and the only quantity of interest is how much downstream capacity fails to recover from it.

Our contribution. We compare eight encoders on a controlled synthetic benchmark designed to make channel identity informative, with supporting experiments on the public ETTh1 dataset (2). Answering the open questions in order: (i) the per-channel linear projection matches every alternative at convergence up to small, statistically real but practically modest, differences—the encoder landscape inside the per-channel- W_k family is one of practical near-equivalence at 20 paired seeds; (ii) the task loss alone drives the learned W_k toward near-orthogonality and scales their norms larger for outcome-relevant channels, with no explicit regulariser required; and (iii) a learned linear projection of the sinusoidal positional encoding gives a small but consistent gain, and a direct geometric probe shows the mechanism is positional-channel orthogonalisation rather than positional subspace compression. The shared-scalar SUM baseline, included as a verified floor, has the capacity-independent ceiling its information-theoretic derivation predicts. Two narrow exceptions to (i) that we resolve under paired analysis but treat as practically small: MLP edges the linear family at $C=16$ with the gap shrinking $\sim 3\times$ under $10\times$ training data, and CAT posts the lowest mean NLL on ETTh1 but is statistically tied with LINEAR and LINEAR-ORTHO.

2 Method

2.1 Model architecture

All encoders share a common backbone: a small causal transformer with $d_{\text{model}}=64$, 4 attention heads, 3 layers, feed-forward width $d_{\text{ff}}=256$, dropout 0.1, pre-LayerNorm, and a linear categorical head mapping to $K=32$ logits. The *only* component that varies across experiments is the input encoder (Figure 1).

2.2 Encoders

We study eight encoders. In the definitions below, $v_k(t)$ is the scalar value of channel k at time t , $\mathbf{p}(t)$ is a fixed sinusoidal positional encoding (1), and all learned parameters are randomly initialised.

sum (ablation baseline). Shared scalar projection $W \in \mathbb{R}^{d_{\text{model}}}$, learned per-channel additive embedding e_k :

$$\mathbf{h}(t) = \sum_k (W v_k(t) + e_k) + \mathbf{p}(t).$$

Total encoder capacity is $\mathcal{O}(d_{\text{model}})$ regardless of C . We use this to isolate what goes wrong when the input projection shares weights across channels.




Figure 1: Architecture. Every experiment uses the same causal transformer backbone (three pre-LN blocks, each structured as in the inset on the right: LN $\rightarrow$ self-attention $\rightarrow$ residual, LN $\rightarrow$ FFN $\rightarrow$ residual; followed by a final LayerNorm and a linear head into $K=32$ bin logits). **The only component that varies across experiments is the encoder (blue).** Eight encoders share the same input/output signature $\mathbb{R}^{B \times T \times C} \rightarrow \mathbb{R}^{B \times T \times d_{\text{model}}}$ and are therefore drop-in interchangeable. The variants are: SUM (shared W then sum across channels), LINEAR (per-channel W_k then sum), LINEAR-ORTHO (LINEAR plus an orthogonality penalty on the W_k), MLP (two-layer MLP on the channel vector), LINEAR-PPE (LINEAR with a learned linear projection of the positional encoding), CONCAT (per-channel block concatenated), CI (channel-independent backbone), and CAT (channel-as-token attention). Precise definitions are in Section 2.

linear. Per-channel projection W_k and bias b_k , summed:

$$\mathbf{h}(t) = \sum_k (W_k v_k(t) + b_k) + \mathbf{p}(t).$$

Equivalently $\mathbf{h}(t) = W \mathbf{v}(t) + \mathbf{b} + \mathbf{p}(t)$ with $W \in \mathbb{R}^{d_{\text{model}} \times C}$: a single `nn.Linear(C, d_model)`.

linear-ortho. Same as LINEAR with an auxiliary loss $L_{\text{ortho}} = \lambda \sum_{i \neq j} (W_i \cdot W_j)^2 / 2$ ($\lambda=10^{-2}$). This variant isolates the regulariser’s contribution.

linear-ppe (projected positional encoding). Same channel encoding as LINEAR, but the sinusoidal position passes through a learned linear layer:

$$\mathbf{h}(t) = \sum_k (W_k v_k(t) + b_k) + W_{\text{pos}} \mathbf{p}(t) + \mathbf{b}_{\text{pos}},$$

with $W_{\text{pos}} \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$. This lets the optimiser rotate the positional subspace to be orthogonal to the channel subspace. Adds $d_{\text{model}}^2 + d_{\text{model}}$ parameters (4,160 at $d=64$).

mlp. Two-layer MLP on the channel vector: $\mathbf{h}(t) = W_2 \text{GELU}(W_1 \mathbf{v}(t) + \mathbf{b}_1) + \mathbf{b}_2 + \mathbf{p}(t)$, with $W_1 \in \mathbb{R}^{d_{\text{model}} \times C}$, $W_2 \in \mathbb{R}^{d_{\text{model}} \times d_{\text{model}}}$. Tests whether nonlinearity helps ($\sim 9\times$ more encoder parameters than LINEAR).

concat. Per-channel projection into $d_{\text{block}}=d_{\text{model}}/C$ dims, concatenated (requires $C \mid d_{\text{model}}$): $\mathbf{h}(t) = [\mathbf{e}_0(t) \parallel \dots \parallel \mathbf{e}_{C-1}(t)] + \mathbf{p}(t)$. Channel identity is an architectural invariant.

ci (channel-independent). PatchTST-spirit (4): shared-weight causal transformer per channel; final hidden states concatenated into the head. Compute: $\mathcal{O}(C B T^2)$.

cat (channel-as-token). iTransformer-spirit (5): each (t, k) pair is a token, sequence length $C \cdot T$, causal across time, bidirectional within a time step. Attention scales as $\mathcal{O}((CT)^2)$, which makes CAT substantially heavier than the other encoders at every C and prohibitively expensive at $C=16, T=160$ on our compute budget; we skip $C=16$ for that reason.

2.3 Training protocol

Identical across encoders and across both datasets unless noted.

Objective. At every position $t \in \{0, \dots, T-1\}$ the model emits logits $\ell_t \in \mathbb{R}^K$ and we minimise the next-step categorical cross-entropy (negative log-likelihood, NLL)

$$\mathcal{L}_{\text{NLL}} = -\frac{1}{B(T-1)} \sum_{b=1}^B \sum_{t=0}^{T-2} \log \text{softmax}(\ell_t^{(b)})_{y_{t+1}^{(b)}},$$

where $y_t^{(b)}$ is the bin index at position t in series b . The predicted distribution at position t is supervised only by the bin at position $t+1$, and causal masking inside the transformer prevents the model from seeing y_{t+1} at prediction time. The LINEAR-ORTHO encoder adds the auxiliary loss $\lambda \sum_{i \neq j} (W_i \cdot W_j)^2 / 2$ with $\lambda=10^{-2}$; all other encoders have $\mathcal{L} = \mathcal{L}_{\text{NLL}}$. As a loss-family robustness check we also re-run a subset of the synthetic sweep with a scalar regression head and MSE loss on the continuous (pre-binning) target; this is described in Section 3.11.

Optimiser. AdamW with peak learning rate 3×10^{-4} , weight decay 10^{-4} , and PyTorch defaults for momentum ($\beta_1=0.9$, $\beta_2=0.999$, $\epsilon=10^{-8}$). Gradient norm is clipped to 1.0 per step. Batch size is 32. No warm-up. Learning rate follows a cosine annealing schedule from peak to 3×10^{-6} (1% of peak) over the full 300 epochs.

Train/val split. For the synthetic benchmark we generate $n_{\text{series}}=512$ length- $T=160$ series per seed and split with a seeded `torch.utils.data.random_split` into 461/51 ($\approx 90/10$) train and val. A separate “data-rich” condition with $n_{\text{series}}=5120$ ($10\times$ training data, same 90/10 split) is used at $C=16$ to check whether the encoder ordering at our main-sweep scale is data-limited; this condition is reported in Section 3.2. For ETTh1 we follow the dataset’s standard chronological partition, using the train block to fit per-variate standardisation statistics and the quantile bin edges for the target, and evaluating on the val block. Data loaders use `shuffle=True` for training, `shuffle=False` for validation, batch size 32 in both cases.

Validation cadence and model selection. Validation NLL and top-1 bin accuracy are evaluated every 20 epochs, plus at epoch 1 and epoch 300 (giving a trace of 17 points). We do not actually halt training early; instead, we report the best validation NLL observed along this trace, together with the accuracy at that same epoch (“effective early stopping”). This decouples our headline metric from slight differences in convergence epoch across encoders while still letting every run see the full 300-epoch schedule.

Seeds and reproducibility. Headline configurations are run over 20 random seeds, indexed 0 through 19. Seed s sets `torch.manual_seed(s)` (controlling weight initialisation and dropout) and seeds the `torch.Generator` passed to `random_split` (controlling the train/val partition); on the synthetic benchmark it also drives the `numpy` RNG that generates the input series. *The same 20 seeds are reused for every encoder variant, every C , and every d_{model} .* This is a paired-seed design: at a fixed s , every encoder sees identical training data and the identical train/val partition, and only the model architecture (and its initial parameter draws from the shared PRNG state) differs. We *report* unpaired mean $\pm$ std intervals throughout for readability, but we *test* with paired-difference statistics where the unpaired interval is loose enough that paired analysis changes the conclusion. Where the paired and unpaired test agree qualitatively (most rows in this paper) we report only the unpaired interval; where they disagree, or where the unpaired p -value is borderline, we report the paired result and call it out. A few diagnostics that consume substantially more compute or are by-design single-instance—linear probing, test-time channel masking, the channel-bias ablation, and the MLP first-layer geometry table—are kept at 5 seeds and labelled as such in their captions. Bootstrap confidence intervals on cosine, paired-difference, and similar derived quantities are constructed by resampling seeds (and, where applicable, examples) with 10,000 resamples; reported as 95% intervals. We do not enable `torch.use_deterministic_algorithms(True)`, so CUDA matmul nondeterminism contributes some additional seed-level variation between full re-runs; the reported $\pm$ std intervals capture this.

2.4 Datasets

Synthetic. $N=512$ time series of length $T=160$ with C channels. Each channel is a sum of three sinusoids with random frequencies in $[0.005, 0.08]$ cycles/sample plus AR(1) noise, standardised. The outcome is a fixed non-linear function of the first four channels with lags 3 and 7:

$$y_t = \tanh(s_0(t-3) s_1(t)) + 0.6 \sin(1.3 s_2(t-7)) + 0.4 \mathbb{I}[s_3(t) > 0] s_0(t),$$

binned into $K=32$ quantile bins. Channels $k \geq 4$ are independent distractors. The task is designed so that an encoder that mixes channels at the input layer cannot recover the interaction structure.

ETTh1. Electricity Transformer Temperature (2), hourly, 7 variates. We standardise on the train split, bin the target variate OT into 32 quantile bins, and predict the next-step bin given all 7 variates over $T=160$. The model uses $d_{\text{model}}=56$ (divisible by $C=7$), 7 heads, $d_{\text{ff}}=224$; training is otherwise identical.

2.5 Diagnostic analyses

Beyond aggregate NLL, we use three diagnostics to inspect what the per-channel- W_k encoders actually learn. Each is computed after full 300-epoch training, on five seeds unless stated otherwise.

Gram-matrix analysis. LINEAR and LINEAR-ORTHO differ only in whether the soft-orthogonality penalty $L_{\text{ortho}} = \lambda \sum_{i \neq j} (W_i \cdot W_j)^2 / 2$ is added to the loss. If both reach the same NLL, the question is whether LINEAR’s W_k are geometrically similar to LINEAR-ORTHO’s—i.e. whether the task pressure alone already produces a near-orthogonal arrangement—or whether the two encoders arrive at the same loss through different geometries.

After training we extract the per-channel projection matrix $W \in \mathbb{R}^{C \times d_{\text{model}}}$ (rows W_k) from the encoder and compute the cosine matrix $G_{ij} = (W_i \cdot W_j) / (\|W_i\| \|W_j\|)$. Two scalar summaries:

$$|\overline{\cos}| = \frac{1}{C(C-1)} \sum_{i \neq j} |G_{ij}|, \quad \max |\cos| = \max_{i \neq j} |G_{ij}|.$$

Both go to zero iff the W_k become mutually orthogonal. We report seed-averaged means.

Encoding-space allocation. Orthogonality says channels are distinguishable; it does not say whether the model dedicates more representational capacity to outcome-relevant channels than to distractors. We measure two quantities, both at the encoder output:

(i) Norm $\|W_k\|$ for each channel. A larger norm means channel k ’s contribution $W_k v_k(t) + b_k$ has larger magnitude in the embedding, i.e. occupies more of the residual stream.

(ii) Per-channel variance fraction. The encoder output is a sum of per-channel contributions $W_k v_k(t) + b_k$. Treating each contribution as a random vector indexed by (t, batch) , the per-channel variance is

$$\sigma_k^2 = \sum_{d=1}^{d_{\text{model}}} \text{Var}_{t, \text{batch}} [(W_k v_k(t) + b_k)_d],$$

and we report fractions $\sigma_k^2 / \sum_j \sigma_j^2$ on a held-out batch. This combines the size of W_k with the empirical variability of channel k ’s input distribution.

For both metrics, channels are partitioned into drivers ($k \in \{0, 1, 2, 3\}$, which enter the outcome formula) and distractors ($k \geq 4$, independent of the outcome). Reporting them this way lets us check whether the model independently discovered the driver/distractor split.

Linear-probing analysis. The encoder lays down a representation; the question is whether downstream attention and FFN layers preserve enough of it that each channel remains linearly identifiable several blocks deep. The pre-LayerNorm transformer’s residual structure already makes a zeroth-order prediction here: each block applies $h \leftarrow h + \text{Attn}(\text{LN}(h))$ and $h \leftarrow h + \text{FFN}(\text{LN}(h))$, so an unmodified copy of the input embedding is carried forward through every block. We should therefore expect deep linear probes to recover whatever was written to the residual stream at layer 0, modulo what LayerNorm rescales and what the sublayers additively perturb. We test this directly by training closed-form ridge probes from a frozen hidden state back to the raw input.

Procedure: (1) train the encoder–transformer end-to-end as usual; (2) freeze it; (3) generate a fresh probe dataset (`MultiSignalDataset` with a different seed, so the model has not been trained on it); (4) for each layer of interest, pass the probe dataset through and collect hidden states $H \in \mathbb{R}^{N \times d_{\text{model}}}$ alongside the corresponding raw inputs $X \in \mathbb{R}^{N \times C}$; (5) solve the ridge regression

$$W_{\text{probe}} = (H^\top H + \lambda I)^{-1} H^\top X, \quad \lambda = 10^{-3},$$

in closed form on a probe-train split; (6) evaluate held-out R^2 per channel on a probe-validation split: $R_k^2 = 1 - \text{SS}_{\text{res},k} / \text{SS}_{\text{tot},k}$.

We probe two layers: layer 0 (the encoder output, i.e. the input to the first transformer block) and layer 3 (the output of the third and final transformer block, before the LayerNorm and the

Table 1: Main comparison at $C=4$, $d_{\text{model}}=64$, 300 epochs, cosine LR decay, mean $\pm$ std over 20 seeds (5 canonical + 15 paired-seed extras). Best val NLL (effective early stopping). Random baseline: $\text{NLL} = \ln 32 \approx 3.47$, $\text{acc}=0.031$. [†]For CI and CAT, “encoder params” counts only the per-channel input projection (plus per-variate embedding for CAT); the bulk of those architectures’ parameters lives in the replacement backbone.

Encoder	Enc. params	Val NLL $\downarrow$	Val acc $\uparrow$
SUM	320	3.257 ± 0.011	0.091 ± 0.003
CI [†]	128	3.053 ± 0.013	0.136 ± 0.004
CAT [†]	384	2.348 ± 0.060	0.212 ± 0.011
CONCAT	128	2.170 ± 0.025	0.243 ± 0.006
LINEAR	512	2.155 ± 0.019	0.247 ± 0.006
LINEAR-ORTHO	512	2.155 ± 0.020	0.247 ± 0.006
MLP	4,480	2.177 ± 0.022	0.242 ± 0.007
LINEAR-PPE	4,672	2.114 ± 0.029	0.256 ± 0.009

categorical head). $R^2 = 1$ at layer 0 confirms the encoder writes the raw channel into a linearly recoverable subspace; R^2 at layer 3 reveals how much of that recoverability survives the depth of the model. CI and CAT are excluded because their hidden states have a fundamentally different shape (per-channel stream and per-token respectively), so a single $H \rightarrow X$ probe does not apply.

3 Results

3.1 Main comparison

Table 1 reports all eight encoders at $C=4$, $d_{\text{model}}=64$, 300 epochs with cosine LR decay, 20 seeds. The headline read is *practical near-equivalence among the per-channel- W_k family*: LINEAR, LINEAR-ORTHO, CONCAT, and MLP all fall in a 0.02 NLL band at 2.155–2.177, with LINEAR and LINEAR-ORTHO matching to three decimals (the orthogonality penalty contributes nothing at convergence). LINEAR-PPE sits a touch below at 2.114, a ~ 0.04 NLL improvement over LINEAR that is small in magnitude but statistically solid under the paired-seed design: 19 of 20 paired differences favour LINEAR-PPE (paired t -test $p=2.6 \times 10^{-6}$, 95% bootstrap CI on Δ NLL $[+0.029, +0.053]$). The mechanism behind the gap is direct: a learned linear projection of the positional encoding rotates it out of the channel subspace (Section 3.10). In practical terms the gap from LINEAR to LINEAR-PPE corresponds to about 2% of the baseline NLL and is well within the spread that small task-design choices would induce; we read it as a real effect, not a practically transformative one. CI and CAT underperform despite nominally giving each channel more capacity. The only encoder that loses decisively is SUM, which lacks per-channel projections.

3.2 Channel-count scaling

All encoders degrade as C grows (channels $k \geq 4$ are distractors). The per-channel- W_k family stays clustered at the top across $C \in \{4, 8, 16\}$ and the gap to SUM widens rather than narrows. LINEAR-PPE sits at the very top across all three C values, with paired tests at 20 seeds: $\Delta = 0.041$, $p=2.6 \times 10^{-6}$ at $C=4$ (19/20 favour ppe); $\Delta = 0.026$, $p=3.7 \times 10^{-4}$ at $C=8$ (16/20); $\Delta = 0.016$, $p=0.036$ at $C=16$ (14/20). The advantage shrinks with C but does not vanish; even at $C=16$ the paired test puts the gap above zero by a small margin.

Table 2: Scaling with channel count. Channels 0–3 drive the outcome; additional channels are independent distractors. 300 epochs with cosine decay, mean $\pm$ std over 20 seeds. CAT is skipped at $C=16$ because its $\mathcal{O}((CT)^2)$ attention cost explodes. Sub-band sub-tables span the within-cell standard error times a t-quantile; encoder-vs-encoder paired tests reported in the text use the same paired-seed design across encoders.

C	encoder	val NLL $\downarrow$	val acc $\uparrow$
4	SUM	3.257 ± 0.011	0.091 ± 0.003
	CI	3.053 ± 0.013	0.136 ± 0.004
	CAT	2.348 ± 0.060	0.212 ± 0.011
	CONCAT	2.170 ± 0.025	0.243 ± 0.006
	MLP	2.177 ± 0.022	0.242 ± 0.007
	LINEAR-ORTHO	2.155 ± 0.020	0.247 ± 0.006
	LINEAR	2.155 ± 0.019	0.247 ± 0.006
	LINEAR-PPE	2.114 ± 0.029	0.256 ± 0.009
8	SUM	3.300 ± 0.006	0.081 ± 0.004
	CI	3.062 ± 0.016	0.135 ± 0.004
	CAT	2.334 ± 0.073	0.213 ± 0.011
	CONCAT	2.234 ± 0.023	0.231 ± 0.005
	MLP	2.211 ± 0.027	0.236 ± 0.005
	LINEAR-ORTHO	2.176 ± 0.022	0.244 ± 0.006
	LINEAR	2.176 ± 0.022	0.243 ± 0.006
	LINEAR-PPE	2.150 ± 0.030	0.249 ± 0.006
16	SUM	3.310 ± 0.004	0.075 ± 0.003
	CI	3.078 ± 0.017	0.133 ± 0.005
	CONCAT	2.340 ± 0.033	0.215 ± 0.008
	LINEAR-ORTHO	2.267 ± 0.038	0.227 ± 0.007
	LINEAR	2.267 ± 0.032	0.226 ± 0.007
	LINEAR-PPE	2.252 ± 0.038	0.230 ± 0.008
	MLP	2.231 ± 0.032	0.234 ± 0.004

At $C=16$, MLP pulls ahead of the linear family (2.231 vs. 2.27–2.34). Paired tests against LINEAR ($p=2 \times 10^{-4}$, 18/20), LINEAR-ORTHO ($p=7 \times 10^{-4}$, 17/20), and CONCAT ($p=5 \times 10^{-12}$, 20/20) are decisive; against LINEAR-PPE the gap is small and the call is borderline ($\Delta = 0.020$, paired $p=0.053$, 12/20). Combined with the LINEAR-PPE lead at smaller C , the picture is a narrow top tier whose member ordering shifts with C : LINEAR-PPE edges the linear family at every C , MLP edges the linear family (but ties LINEAR-PPE) at $C=16$, and at every C the absolute spread among the top four is on the order of 0.02–0.05 NLL on a baseline of ~ 2.2 NLL.

The $C=16$ ranking persists, but shrinks, with $10\times$ training data. The $\Delta\text{NLL} \sim 0.07$ between MLP and LINEAR at $C=16$, $N_{\text{series}}=512$ raises an obvious question: is the nonlinear stem genuinely useful, or is the main sweep data-limited and MLP just happens to extract more from sparse data? We test this directly with the `main.largen` stage of our reproducer: the same top-tier encoders at $C=16$ trained with $N_{\text{series}}=5120$ ($10\times$ the main-sweep data), 20 seeds, otherwise identical.

Two things happen at once. (i) MLP’s lead survives in the data-rich regime: 18–20 of 20 paired-seed differences favour MLP over each of LINEAR (paired $p=4 \times 10^{-4}$), LINEAR-ORTHO ($p=9.6 \times 10^{-4}$),




Figure 2: Scaling with input channel count C (error bars: one seed std over 20 seeds). CAT skipped at $C=16$ due to $\mathcal{O}((CT)^2)$ attention cost.

Table 3: $C=16$ top-tier encoders, $N_{\text{series}}=5120$ vs. $N_{\text{series}}=512$ (20 seeds each, 300 epochs). The full linear family at $N=5120$ falls inside a 0.024 NLL band.

encoder	$N=512$ val NLL	$N=5120$ val NLL	Δ from data
MLP	2.231 ± 0.032	1.831 ± 0.009	-0.400
LINEAR-PPE	2.252 ± 0.038	1.841 ± 0.013	-0.411
LINEAR	2.267 ± 0.032	1.843 ± 0.015	-0.424
LINEAR-ORTHO	2.267 ± 0.038	1.844 ± 0.018	-0.423
CONCAT	2.340 ± 0.033	1.854 ± 0.011	-0.486

LINEAR-PPE ($p=2.4 \times 10^{-4}$), and CONCAT ($p=2.3 \times 10^{-13}$). The $C=16$ MLP finding is not a data-limited artefact. (ii) The magnitude shrinks roughly three-fold: $\Delta\text{NLL}(\text{MLP-LINEAR})$ goes from 0.036 at $N=512$ to 0.011 at $N=5120$; the full linear family at $N=5120$ falls inside a band of width 0.024 NLL versus 0.073 at $N=512$. Each encoder also gains ~ 0.4 NLL from the extra data, confirming the $N=512$ sweep is data-limited at $C=16$ just as at $C=4$ (Section 3.6). The encoder ordering at $C=16$ is robust to data scale; the absolute size of the gaps is not.

3.3 Model-width scaling

Table 4 varies d_{model} over $\{64, 128, 256\}$ at $C=4$, 20 seeds, all top-tier encoders plus the SUM and CONCAT baselines.

Three findings. First, SUM’s ceiling is structural and capacity-independent: $15\times$ more parameters move NLL by less than 0.01, as expected from the information-theoretic argument later in Section 3.5.

Second, *every* encoder shows clear overfitting at $d_{\text{model}} \geq 128$. Best val NLL is reached at epoch ~ 100 ($d=128$) or ~ 40 ($d=256$), after which val NLL drifts upward (catastrophically for MLP at $d=256$, where the end-of-training val NLL exceeds $\ln 32$). Effective early stopping (Section 2.3) is what protects the reported best-NLL numbers; the trajectories themselves are not converged. $C=4$ with $N_{\text{series}}=512$ does not justify a 256-dim residual stream.

Third, the linear family’s best-val ceiling rises gently with d_{model} ($2.16 \rightarrow 2.23 \rightarrow 2.32$ for LINEAR), indicating that the optimal effective-stopping model gets worse as the model is given more

Table 4: Scaling with d_{model} at $C=4$. $d_{\text{ff}}=4d_{\text{model}}$, 4 attention heads (so d_{head} grows). Mean $\pm$ std over 20 seeds, 300 epochs with cosine decay. *All* encoders flag as overfitting at $d_{\text{model}} \geq 128$: best val NLL is reached well before epoch 300 (typical best epochs ~ 100 at $d=128$, ~ 40 at $d=256$), and val NLL drifts upward afterwards. The reported best-NLL numbers are the effective-early-stopping minima.

d_{model}	encoder	total params	val NLL $\downarrow$	val acc $\uparrow$
64	SUM	152,480	3.257 ± 0.011	0.091 ± 0.003
	CONCAT	152,288	2.170 ± 0.026	0.243 ± 0.006
	LINEAR	152,672	2.155 ± 0.019	0.248 ± 0.005
	LINEAR-ORTHO	152,672	2.155 ± 0.019	0.248 ± 0.005
	MLP	156,640	2.177 ± 0.023	0.241 ± 0.008
	LINEAR-PPE	156,832	2.115 ± 0.029	0.256 ± 0.009
128	SUM	599,840	3.260 ± 0.010	0.091 ± 0.003
	CONCAT	599,456	2.229 ± 0.031	0.235 ± 0.006
	LINEAR	600,224	2.234 ± 0.032	0.235 ± 0.006
	LINEAR-ORTHO	600,224	2.233 ± 0.032	0.235 ± 0.007
	MLP	616,352	2.301 ± 0.027	0.222 ± 0.007
	LINEAR-PPE	616,736	2.185 ± 0.032	0.245 ± 0.005
256	SUM	2,379,296	3.262 ± 0.011	0.091 ± 0.003
	CONCAT	2,378,528	2.311 ± 0.026	0.217 ± 0.004
	LINEAR	2,380,064	2.322 ± 0.033	0.217 ± 0.006
	LINEAR-ORTHO	2,380,064	2.323 ± 0.036	0.217 ± 0.006
	MLP	2,445,088	2.407 ± 0.023	0.210 ± 0.007
	LINEAR-PPE	2,445,856	2.265 ± 0.026	0.229 ± 0.006

capacity to overfit. MLP degrades fastest of the top tier ($2.18 \rightarrow 2.30 \rightarrow 2.41$) because its encoder parameters scale as d_{model}^2 and exacerbate the capacity mismatch. The paired tests against LINEAR sharpen with width: $p=2 \times 10^{-4}$ at $d=64$ (18/20 favour LINEAR), $p=4 \times 10^{-9}$ at $d=128$ (20/20), and $p=5 \times 10^{-9}$ at $d=256$ (19/20). The $C=4$ small-channel regime does not reward nonlinearity, and gives it room to overfit when widened. This is in sharp contrast to $C=16$ where MLP leads (Section 3.2): the high-channel-pressure regime is what makes nonlinearity useful.

linear-ppe continues to lead. At every d_{model} tested, LINEAR-PPE achieves the lowest val NLL. The paired gaps to LINEAR are $\Delta=0.040$ at $d=64$ (paired $p=3 \times 10^{-6}$, 19/20), $\Delta=0.049$ at $d=128$ (paired $p=3 \times 10^{-4}$, 18/20), and $\Delta=0.056$ at $d=256$ (paired $p=4 \times 10^{-8}$, 19/20). The advantage is robust (56/60 paired-seed comparisons across the three widths favour LINEAR-PPE), and the magnitude grows mildly with d_{model} . Both observations are consistent with the direct geometric measurement in Section 3.10: the rotation is doing real work, and more residual-stream capacity gives it more directions to rotate into.

3.4 Gram-matrix analysis: spontaneous orthogonality

LINEAR has no mechanism pushing the W_k toward orthogonality; LINEAR-ORTHO adds an explicit penalty. Both produce the same NLL to three decimals. Applying the off-diagonal cosine diagnostic from Section 2.5 gives:

Table 5: Learned per-channel W_k geometry ($C=4$, 300 epochs, mean over 20 seeds).

encoder	best val NLL	$\max_{i \neq j} \cos(W_i, W_j) $	mean $ \cos $
LINEAR	2.155	0.095	0.042
LINEAR-ORTHO	2.155	0.022	0.010

Without any penalty the learned W_k are already near-orthogonal (mean off-diagonal cosine 0.042, 95% bootstrap CI [0.037, 0.046]). The regulariser tightens this by another factor of ~ 4 to 0.010 (95% CI [0.008, 0.011]) but does not improve downstream NLL. The task loss itself provides enough gradient pressure to separate the per-channel projections once they exist as free parameters.

Does the MLP stem do the same thing? The same diagnostic applies to MLP’s first-layer weight $W^{(1)} \in \mathbb{R}^{d_{\text{hidden}} \times C}$: its columns are the per-channel input directions before the GELU nonlinearity, and their pairwise cosines measure how much the encoder separates channels at the linear-stem stage versus deferring that separation to GELU plus $W^{(2)}$. Training MLP at $C \in \{4, 8, 16\}$ with 5 seeds and applying the off-diagonal cosine diagnostic to $W^{(1)}$ ’s columns:

Table 6: MLP first-layer column geometry (300 epochs, $d_{\text{model}}=64$, 5 seeds). Compare to LINEAR’s W_k at $C=4$: mean 0.042, max 0.095.

C	best val NLL	mean $ \cos $	max $ \cos $
4	2.170 ± 0.036	0.075 ± 0.021	0.158 ± 0.058
8	2.214 ± 0.016	0.074 ± 0.015	0.248 ± 0.098
16	2.222 ± 0.025	0.105 ± 0.005	0.463 ± 0.109

MLP’s linear stem is also pushed toward near-orthogonality by the task loss, but consistently looser than the linear family’s W_k : at $C=4$ its mean off-diagonal cosine is $1.8\times$ LINEAR’s (0.075 vs. 0.042); at $C=16$ the gap widens further and pairs of columns reach $|\cos| \approx 0.46$. The mechanism is straightforward: MLP can also separate channels through the GELU plus $W^{(2)}$, so the gradient pressure on $W^{(1)}$ alone is partially relieved. The qualitative finding survives—spontaneous channel orthogonalisation is a property of the task loss interacting with a per-channel input stem, not unique to a single-layer projection—but quantitatively the linear family achieves the cleanest geometry.

3.5 Bias ablation: the biases play no role in the geometry

LINEAR’s forward pass is $h(t) = \sum_k (W_k v_k(t) + b_k) + \mathbf{p}(t)$. Distributing the sum, the per-channel biases collapse into a single d_{model} -dimensional constant offset $B = \sum_k b_k$ that is added to every embedding regardless of input or position. Algebraically, the $C \times d_{\text{model}}$ bias parameters are an over-parameterised reparameterisation of one global offset, and they never enter the cosine computation in Table 5. To confirm experimentally that they play no role in the orthogonalisation, we retrain LINEAR with the `channel_bias` parameter zeroed and frozen (LINEAR-NOBIAS):

NLL is unchanged within seed std, and the Gram statistics match to three decimals. The $C \cdot d_{\text{model}} = 256$ bias parameters at this configuration are 0.17% of the full model’s $\sim 152\text{K}$ parameters, so the saving is too small to drive a practical recommendation; we retain the bias in the headline definition of LINEAR to match the PyTorch `nn.Linear` default. The point of this ablation is epistemic: it confirms that the geometric argument in this paper is bias-independent.

Table 7: Channel-bias ablation ($C=4$, 300 epochs, 5 seeds).

encoder	val NLL ↓	mean cos	max cos
LINEAR	2.169 ± 0.014	0.036	0.082
LINEAR-NOBIAS	2.177 ± 0.013	0.036	0.083

3.6 Encoding space allocation

Beyond separating channels, does the model allocate more encoding space to channels that matter more? Using the norm and variance-fraction diagnostics from Section 2.5, at $C \in \{4, 8, 16\}$:

Table 8: Encoding space allocation for LINEAR (300 epochs, mean over 20 seeds). Channels 0–3 drive the outcome; the rest are distractors.

	$\ W_k\ $ norm		variance fraction	
	drivers (0–3)	distractors	drivers	distractors
$C=8$	1.22	0.54	83.8%	16.2%
$C=16$	1.30	0.59	62.4%	37.6%

Driver channels receive $\sim 2.2\times$ larger seed-averaged norms than distractors at $C=8$ and $\sim 2.1\times$ at $C=16$. The four drivers capture 83% of the embedding variance at $C=8$ and 61% at $C=16$, where the aggregate share is naturally diluted by having twelve distractors but the per-channel variance contribution remains markedly larger for drivers ($\sim 15\%$ each vs. $\sim 3\%$ each). We do not report per-channel error bars on the norms, and the fine-grained ordering within drivers is descriptive only; the robust claim is the driver/distractor gap, which holds at both channel counts.

Why distractor norms are not zero. The driver/distractor gap is large but the distractor norms (~ 0.55 at $C=8$) do not vanish. This is a finite-data equilibrium rather than an asymptotic property of the encoder. The expected gradient on a distractor weight W_k is

$$\mathbb{E}\left[\frac{\partial \mathcal{L}}{\partial W_k}\right] = \mathbb{E}\left[\frac{\partial \mathcal{L}}{\partial \mathbf{h}(t)} \cdot v_k(t)\right],$$

which factorises to zero whenever the distractor input v_k is independent of the drivers and the outcome (true by construction). With finite data, however, stochastic gradient estimates have nonzero variance scaling as $\mathcal{O}(1/\sqrt{N_{\text{series}}})$, and under L2 weight decay ($\lambda=10^{-4}$) distractor weights sit at a noise floor of order $\sqrt{\sigma_{\text{noise}}^2/\lambda}$ rather than being driven to zero. The $1/\sqrt{N}$ scaling predicts that the $N=5120$ data-rich condition (Section 3.2) should shrink distractor norms by $\sim \sqrt{10} \approx 3.2$; running LINEAR at $C=8$ in that condition gives a $\sim 2.5\times$ reduction (driver/distractor ratio rises from $2.2\times$ to $6.7\times$, driver variance share rises from 83% to 98%), close to but slightly slower than predicted. Sparser regularisers (L_1 on rows of W , larger weight decay) would push the distractor norms to exact zero, but at the cost of deviating from the bare `nn.Linear(C,  $d_{\text{model}}$ )` we are auditing; we do not employ them because the qualitative result—“the model down-weights distractors and sharpens that down-weighting with more data”—is already supported by the norm gap, the variance split, and the N -scaling probe.

Table 9: Linear-probe channel recovery at $C=4$, seed 0, 300 epochs. Closed-form ridge probe $\mathbf{h} \mapsto \hat{x}$ from a frozen hidden state to the raw input $x \in \mathbb{R}^4$; held-out per-channel R^2 . Layer 0 is the input embedding; layer 3 is the final transformer block. All per-channel- W_k encoders achieve near-perfect input recovery and preserve mean $R^2 \geq 0.84$ through three attention layers; SUM collapses to $R^2 \approx 0.24$ at every depth.

Encoder	Layer	R^2_{ch0}	R^2_{ch1}	R^2_{ch2}	R^2_{ch3}	mean
SUM	0 (input)	0.255	0.250	0.242	0.214	0.240
SUM	3 (final)	0.248	0.242	0.235	0.206	0.233
LINEAR	0 (input)	1.000	1.000	1.000	1.000	1.000
LINEAR	3 (final)	0.957	0.861	0.911	0.879	0.902
LINEAR-ORTHO	0 (input)	1.000	1.000	1.000	1.000	1.000
LINEAR-ORTHO	3 (final)	0.959	0.856	0.915	0.885	0.904
MLP	0 (input)	1.000	1.000	1.000	1.000	1.000
MLP	3 (final)	0.934	0.879	0.913	0.646	0.843
LINEAR-PPE	0 (input)	1.000	1.000	1.000	1.000	1.000
LINEAR-PPE	3 (final)	0.930	0.861	0.874	0.847	0.878
CONCAT	0 (input)	0.998	1.000	1.000	1.000	0.999
CONCAT	3 (final)	0.955	0.852	0.906	0.845	0.890

3.7 Channel identifiability via linear probing

Following the procedure in Section 2.5, all per-channel- W_k encoders achieve $R^2 \approx 1.0$ at the encoder output (layer 0) and $R^2 \geq 0.84$ at the output of the third transformer block. SUM collapses to $R^2 \approx 0.24$ at every depth—channel identity is unrecoverable from the embedding it produces.

3.8 Channel masking

The outcome formula uses ch0 in two interaction terms, ch1 in one, and ch2/ch3 once each with smaller marginal effect—so we know a priori which channels the model *should* rely on most. Zeroing one channel at test time gives us a behavioural test of whether the encoder reproduces this ordering. Table 10: the top-tier encoders show sharp, interpretable per-channel accuracy drops that match the outcome formula. SUM is nearly flat, as expected from the information-theoretic argument: it has no recoverable channel identity for the masking to disturb.

3.9 Convergence and wall-clock cost

Two complementary worries about more involved encoders: (i) that they cost extra training time even when they don’t change the ceiling, and (ii) the converse—that they might, despite tying or losing on best-NLL, at least *reach* that ceiling faster and therefore be the right choice when compute or epochs are constrained. From the 300-epoch traces (val NLL recorded every 20 epochs, $C=4$, 5 seeds) we extract two diagnostics: the epoch at which each run first reaches within 0.05 NLL of its own best, and the wall-clock seconds per epoch, normalised to LINEAR.

The five top-tier encoders reach within 0.05 NLL of their final value at the same trace point (~ 150 epochs) and cost essentially the same per-epoch wall time—the transformer backbone dominates, so encoder choice within this family is free. Neither worry materialises: the more involved encoders do not slow convergence (the MLP stem’s $9\times$ extra parameters cost nothing per epoch and reach target on the same trace point), but they also do not converge faster than LINEAR. There is no

Table 10: Test-time channel ablation at $C=4$, seed 0, 300 epochs. We zero a single channel’s input at inference; rows show the resulting val accuracy. Δ is accuracy drop from the unmasked baseline. The outcome formula uses ch0 in two interaction terms, ch1 in one, and ch2/ch3 once each with smaller marginal effect; the per-channel- W_k encoders reproduce this ordering cleanly, whereas SUM, CI, and CAT are flatter.

encoder	no mask	mask ch0	mask ch1	mask ch2	mask ch3
SUM	0.093	0.082	0.089	0.084	0.094
Δ acc	—	−0.011	−0.004	−0.009	+0.001
CI	0.144	0.082	0.108	0.086	0.134
Δ acc	—	−0.062	−0.036	−0.058	−0.010
CAT	0.217	0.095	0.110	0.110	0.200
Δ acc	—	−0.122	−0.107	−0.107	−0.017
CONCAT	0.241	0.096	0.113	0.103	0.202
Δ acc	—	−0.145	−0.128	−0.138	−0.039
LINEAR	0.238	0.096	0.116	0.106	0.204
Δ acc	—	−0.142	−0.122	−0.132	−0.034
LINEAR-ORTHO	0.239	0.096	0.117	0.104	0.205
Δ acc	—	−0.143	−0.122	−0.135	−0.034
MLP	0.240	0.098	0.111	0.106	0.209
Δ acc	—	−0.142	−0.129	−0.134	−0.031
LINEAR-PPE	0.256	0.096	0.115	0.106	0.216
Δ acc	—	−0.160	−0.141	−0.150	−0.040

“faster-to-target” argument available for picking MLP, LINEAR-ORTHO, or LINEAR-PPE over LINEAR within the top tier.

CI and CAT pay substantial wall-clock costs from their architectures, not their input encoders: $\sim 2.3\times$ LINEAR at $C=4$ for CI (scaling to $\sim 8\times$ at $C=16$) and $\sim 5\times$ at $C=4$ for CAT (scaling with C^2T^2 to $\sim 17\times$ at $C=8$). On the synthetic benchmark CAT also converges to its worse ceiling later than the top tier reaches its ceiling.

3.10 Positional projection: a small but real effect across C

LINEAR-PPE achieves the lowest val NLL at every C we test. The paired-seed design (Section 2.3) gives one Δ_s per seed; at 20 seeds the picture is:

- $C=4$: $\Delta = 0.041$, paired $t=6.60$, $p=2.6\times 10^{-6}$, 95% bootstrap CI $[+0.029, +0.053]$, 19/20 favour ppe.
- $C=8$: $\Delta = 0.026$, paired $p=3.7\times 10^{-4}$, 95% bootstrap CI $[+0.014, +0.038]$, 16/20.
- $C=16$: $\Delta = 0.016$, paired $p=0.036$, 95% bootstrap CI $[+0.002, +0.028]$, 14/20.

The gap shrinks with C but does not vanish; even at $C=16$, where the earlier 5-seed analysis read as null, the 20-seed paired test puts Δ above zero, just barely. The advantage is small in magnitude— $\sim 2\%$ of baseline NLL at $C=4$, $\sim 0.7\%$ at $C=16$ —and the practical-significance reading is more subdued than the statistical-significance one would suggest in isolation.

Table 11: Convergence speed and wall-clock cost ($C=4$, 300 epochs, 5 seeds, single GPU). Epochs-to-target is rounded to the trace resolution of 20 epochs. Cost is per-epoch wall time normalised to LINEAR.

encoder	best NLL	epoch to NLL +0.05	relative cost
LINEAR	2.170	~ 152	$1.00\times$
LINEAR-ORTHO	2.170	~ 156	$1.04\times$
MLP	2.171	~ 144	$1.01\times$
CONCAT	2.183	~ 152	$1.00\times$
LINEAR-PPE	2.116	~ 148	$0.99\times$
SUM	3.252	~ 20 (plateau)	$1.00\times$
CI	3.054	~ 20 (plateau)	$2.3\times$
CAT	2.360	~ 180	$5.2\times$

Direct measurement: orthogonalisation, not compression. There are two natural mechanisms by which a learned linear projection on top of $\mathbf{p}(t)$ could help. The *orthogonalisation* reading: W_{pos} rotates the positional basis out of $\text{span}(W)$, so position and channels stop competing for the same coordinates of the residual stream. The *compression* reading: W_{pos} projects $\mathbf{p}(t)$ onto a small effective subspace, leaving the remaining dimensions free for channels. The two predictions differ: orthogonalisation should reduce the overlap between the positional basis and $\text{span}(W)$ while leaving the basis’s intrinsic rank essentially unchanged; compression should reduce that rank. We test both directly.

We extract the per-channel projection matrix $W \in \mathbb{R}^{C \times d_{\text{model}}}$ and the effective positional basis $P \in \mathbb{R}^{T \times d_{\text{model}}}$ from trained LINEAR and LINEAR-PPE models ($C=4$, 20 seeds each). For LINEAR, $P_{t,:} = \mathbf{p}(t)$ is the fixed sinusoidal basis; for LINEAR-PPE, $P_{t,:} = W_{\text{pos}}\mathbf{p}(t) + \mathbf{b}_{\text{pos}}$ is the learned-rotated basis. Two metrics suffice to discriminate the mechanism:

Table 12: Geometric measurement of the positional basis P and its overlap with the channel subspace $\text{span}(W)$ ($C=4$, 300 epochs, mean over 20 seeds with 95% bootstrap CIs in brackets where seed variance is non-zero).

	LINEAR	LINEAR-PPE
Effective rank of P (entropy of $\tilde{\sigma}^2$)	7.59	9.55 [9.17, 9.97]
Fraction of $\ P\ _F^2$ within $\text{span}(W)$	0.034 [0.029, 0.039]	0.005 [0.005, 0.006]

The two metrics give opposite verdicts on the two readings. Orthogonalisation is confirmed: the fraction of P ’s energy that lies inside $\text{span}(W)$ drops by $\sim 6.3\times$ under LINEAR-PPE ($3.4\% \rightarrow 0.5\%$; paired-difference 95% bootstrap CI on the reduction $[-0.034, -0.024]$), so the learned W_{pos} rotation actively pushes the positional encoding out of the channel subspace. Compression is contradicted: P ’s effective rank is *higher* under LINEAR-PPE (9.55 vs. 7.59), not lower—the rotation spreads positional energy across more directions, not fewer. The mechanism is rotation, not compression.

The redistribution is visible at the level of individual singular values: LINEAR’s P has one dominant direction ($\sigma_0 \approx 52$) plus a long tail; LINEAR-PPE’s is more evenly distributed ($\sigma_0 \approx 25$, with the next several around 11–13). Rotating the dominant positional mode out of $\text{span}(W)$ inevitably redistributes its energy across other directions, which is the same effect read either as “orthogonalising position from channels” or as “flattening the singular spectrum of P ”.

Table 13: Loss-family robustness check: MSE/regression target replaces categorical cross-entropy. Same backbone, 5 seeds, 300 epochs, $d_{\text{model}}=64$. Lower MSE is better. The ordering survives the swap: SUM collapses ($R^2 \approx 0.13$ at $C=4$, ≈ 0.01 at $C=16$), the per-channel- W_k tier (CONCAT, LINEAR, LINEAR-ORTHO, MLP, LINEAR-PPE) clusters tight, and LINEAR-PPE edges LINEAR at $C=4$ ($\sim 6\%$ MSE reduction). MLP pulls ahead of LINEAR at $C=16$, mirroring the categorical sweep.

encoder	$C=4$		$C=16$	
	val MSE $\downarrow$	val R^2 $\uparrow$	val MSE $\downarrow$	val R^2 $\uparrow$
SUM	0.4719 ± 0.0090	0.128 ± 0.016	0.5338 ± 0.0167	0.011 ± 0.006
CONCAT	0.0636 ± 0.0027	0.882 ± 0.006	0.0862 ± 0.0071	0.840 ± 0.014
LINEAR	0.0619 ± 0.0024	0.886 ± 0.005	0.0736 ± 0.0026	0.863 ± 0.008
LINEAR-ORTHO	0.0619 ± 0.0024	0.886 ± 0.005	0.0735 ± 0.0020	0.864 ± 0.007
MLP	0.0643 ± 0.0044	0.881 ± 0.009	0.0678 ± 0.0025	0.874 ± 0.007
LINEAR-PPE	0.0582 ± 0.0031	0.892 ± 0.007	0.0722 ± 0.0033	0.866 ± 0.007

We also computed a third candidate diagnostic—principal angles between $\text{span}(W)$ and $\text{span}(P)$ —but found it degenerate at this dimensionality: $\text{span}(W)$ is only 4-dimensional, $\text{span}(P)$ has rank ~ 22 in a 64-dimensional ambient space, and $\text{span}(W)$ is therefore trivially contained in $\text{span}(P)$ for either encoder (all principal angles come out ~ 0). Principal angles measure subset membership, not energetic overlap; the fraction-of-energy metric is the right tool here.

With orthogonalisation confirmed, the C -dependent shrinkage of the gap admits a natural mechanistic reading: W_{pos} must rotate P out of a C -dimensional channel span using a fixed d_{model}^2 parameter budget, and the rotation becomes increasingly over-constrained as C grows. At $C=4$ ($d_{\text{model}}-C = 60$ “free” dimensions for P to occupy) the rotation clearly helps; at $C=16$ ($d_{\text{model}}-C = 48$, plus more channel pressure on the residual stream overall) it helps less. We cannot fully separate this from two other plausible factors—the synthetic task has only four drivers regardless of C , so representational effort shifts elsewhere at large C ; and seed-level optimisation variance is larger at higher C —but the over-constraint reading is the only one that directly invokes the empirically-supported mechanism.

The LINEAR-PPE lead also extends across model widths: see the d_{model} sweep in Section 3.3, where LINEAR-PPE achieves the lowest val NLL at every $d_{\text{model}} \in \{64, 128, 256\}$.

3.11 Loss-family robustness check: MSE target

The categorical $K=32$ -bin head could in principle interact with encoder choice—e.g. by rewarding encoders that resolve fine target quantiles—so we rerun a subset of the synthetic main sweep with a scalar regression head trained on the continuous (pre-binning) target under MSE loss. Same backbone, same training schedule, 5 seeds, $C \in \{4, 16\}$, top-tier encoders plus the SUM baseline (Table 13). The ordering is preserved at both channel counts: SUM collapses to $R^2 \approx 0.13$ at $C=4$ and ≈ 0.01 at $C=16$, consistent with the information-theoretic ceiling argument; the per-channel- W_k tier clusters tight; LINEAR-PPE edges LINEAR at $C=4$ ($\sim 6\%$ MSE reduction), and MLP edges LINEAR at $C=16$. The categorical-vs.-MSE choice is therefore not what is driving the encoder ordering—both losses penalise SUM for the same information-theoretic reason, and both reveal the same small gaps inside the top tier. The MSE sweep is not powered (5 seeds, no paired design) to resolve the within-tier gaps at the statistical level of the categorical 20-seed analysis; what it does is rule out that the entire ordering is a label-binning artefact.

Table 14: ETTh1 next-step bin prediction. 7 variates, $T=160$, $K=32$ quantile bins on the target OT (oil temperature) computed from the train split. Target is the next-step bin of OT; all 7 variates are inputs. $d_{\text{model}}=56$, 7 heads, 3 layers, $d_{\text{ff}}=224$; 300 epochs, cosine decay, mean $\pm$ std over 20 seeds (5 canonical + 15 paired-seed extras).

encoder	val NLL $\downarrow$	val acc $\uparrow$
SUM	3.668 ± 0.063	0.016 ± 0.020
CI	0.865 ± 0.038	0.664 ± 0.018
CAT	0.551 ± 0.019	0.785 ± 0.010
LINEAR	0.561 ± 0.017	0.786 ± 0.008
LINEAR-ORTHO	0.561 ± 0.017	0.784 ± 0.009
CONCAT	0.571 ± 0.019	0.783 ± 0.013
LINEAR-PPE	0.573 ± 0.014	0.776 ± 0.009
MLP	0.585 ± 0.020	0.788 ± 0.011

3.12 Real-data validation: ETTh1

Table 14 mostly confirms the synthetic ordering on ETTh1: the per-channel- W_k family (LINEAR, LINEAR-ORTHO, CONCAT, LINEAR-PPE, MLP) cluster within one standard deviation across seeds of one another, CI underperforms, and SUM fails catastrophically (NLL near $\ln 32 \approx 3.47$).

One difference from the synthetic ranking: CAT, which sits in the middle of the synthetic table, here posts the lowest mean NLL (0.551). Paired-difference tests at 20 seeds give a nuanced picture:

- Decisive: CAT beats MLP (paired $p=1.4 \times 10^{-7}$, 20/20) and LINEAR-PPE ($p=3 \times 10^{-4}$, 17/20).
- Marginal: CAT beats CONCAT ($p=0.002$, 16/20).
- Indistinguishable: CAT vs. LINEAR ($\Delta = 0.010$, paired $p=0.14$, 13/20) and CAT vs. LINEAR-ORTHO ($\Delta = 0.010$, paired $p=0.10$, 13/20).

Best-bin accuracies across the top tier are indistinguishable (0.78–0.79 with standard deviation across seeds ~ 0.01). So CAT posts the lowest mean NLL but is statistically tied with LINEAR and LINEAR-ORTHO at the seed-level we test. Read carefully, the finding is not that channel-as-token decisively wins on real data, but that its synthetic-benchmark disadvantage closes and a small directional lead emerges—one that does not survive a paired test against the closest competitors. A plausible reading is that ETTh1’s seven variates are all genuinely informative measurements of the same physical system, so the fine-grained cross-variate attention CAT enables can pay; whether that pays *enough* to prefer it over the per-channel- W_k default at matched wall-clock budget is much less clear.

4 Discussion

linear is not new. Stacking per-channel projections W_k and summing is `nn.Linear(C,  $d_{\text{model}}$ )`—the most obvious default. Our contribution is not proposing it but auditing it: at matched parameter budget, it matches every alternative (block-partitioned, orthogonality-regularised, nonlinear, channel-independent, channel-as-token) on the synthetic benchmark, and ties at the top of the per-channel- W_k tier on ETTh1. Only the shared-projection SUM consistently loses.

What fails about the shared-scalar baseline. SUM collapses because, with a shared scalar projection W and constant per-channel embeddings, the encoder depends on its inputs only through their pointwise sum $S(t) = \sum_k v_k(t)$. That map has a $(C-1)$ -dimensional null space, and by the data-processing inequality no downstream layer can recover what is destroyed before the transformer sees it. The probe makes the floor explicit: for independent unit-variance channels at $C=4$ the theoretical ceiling for linearly recovering one channel from S is $\rho^2 = 1/C = 0.25$, and the layer-0 probe of SUM sits at $R^2 \approx 0.24$, touching it (Table 9). The $15 \times d_{\text{model}}$ sweep in Table 4 moves NLL by less than 0.005 because the extra parameters never get to see the lost information. The bottleneck is not summation per se—LINEAR also sums—but the *shared* projection, which collapses the per-channel directions into one before the sum.

Orthogonality is spontaneous. LINEAR’s W_k reach mean off-diagonal cosine ≈ 0.042 with no penalty; the regulariser tightens to ≈ 0.010 without changing NLL. The task loss provides the gradient pressure.

Importance weighting is automatic. At $C=8$, driver channels get $\sim 2.3 \times$ larger seed-averaged $\|W_k\|$ norms and 84% of the embedding variance; at $C=16$ the norm ratio is $\sim 2.2 \times$ and per-channel variance contributions are $\sim 4 \times$ larger for drivers than for distractors. The driver/distractor gap is robust across two channel counts; the fine-grained ordering within drivers is descriptive only.

The residual stream carries channel identity end-to-end. Linear probes at layer 3 recover each channel to $R^2 \geq 0.84$, confirming the residual-stream prediction in Section 2.5: per-channel- W_k encoders write near-orthogonal channel directions into the residual stream at layer 0 and the pre-LN architecture preserves them through every block essentially by construction, since LayerNorm only rescales and the additive residual keeps them in the basis. The geometric argument in this paper is self-consistent: the encoder chooses the subspace, the backbone passes it forward, and the deep-probe recoverability falls out of both.

When nonlinearity helps and when it hurts. MLP is statistically worse than LINEAR at $C=4, d_{\text{model}}=64$ by a small margin ($\Delta = 0.022$, paired $p=2 \times 10^{-4}$, 18/20) despite its $9 \times$ more encoder parameters; widening the model at $C=4$ makes things worse still ($\Delta = 0.067$ at $d=128$, $\Delta = 0.085$ at $d=256$, both paired $p < 10^{-8}$). MLP’s encoder parameters scale as d_{model}^2 , which exacerbates the overfitting that the $N=512$ regime already produces at $d \geq 128$ (Section 3.3). At $C=16$ the picture flips: paired tests (Section 3.2) show MLP ahead of LINEAR ($\Delta = 0.036$, $p=2 \times 10^{-4}$, 18/20), LINEAR-ORTHO ($\Delta = 0.036$, $p=7 \times 10^{-4}$), and CONCAT ($\Delta = 0.108$, $p=5 \times 10^{-12}$); against LINEAR-PPE the gap is small and borderline ($\Delta = 0.020$, paired $p=0.053$, 12/20). A plausible reading: with many channels (drivers plus more distractors) the nonlinearity buys the encoder a useful gating on the channel vector that a linear projection cannot match. At $10 \times$ training data (Table 3) MLP’s lead persists (18/20 to 20/20 paired diffs favouring MLP, paired $p \in [2 \times 10^{-13}, 4 \times 10^{-4}]$) but shrinks $\sim 3 \times$ to $\Delta = 0.010$ – 0.023 NLL, indicating that the $N=512$ gap is partly amplified by the data-limited regime. The gap is small, shrinks further with more data, and the architecture is more complex; on this evidence we do not recommend MLP as a default choice even at the $C=16$ regime where it leads.

Projected positional encoding: small but real across C . The paired-seed analysis in Section 3.10 establishes the LINEAR-PPE gap over LINEAR as small but statistically clear at every C tested: $\Delta = 0.041$ ($p=2.6 \times 10^{-6}$) at $C=4$, $\Delta = 0.026$ ($p=3.7 \times 10^{-4}$) at $C=8$, $\Delta = 0.016$ ($p=0.036$)

at $C=16$. The gap shrinks with C but does not vanish. A direct geometric measurement (Table 12) settles the mechanism: the fraction of the positional basis P ’s energy that lies inside the channel subspace $\text{span}(W)$ drops $\sim 6.3\times$ under LINEAR-PPE ($3.4\% \rightarrow 0.5\%$); the learned W_{pos} rotation actively pushes P out of $\text{span}(W)$. The orthogonalisation reading is confirmed; the compression reading is contradicted by the same measurement (P ’s effective rank is *higher* under LINEAR-PPE, 9.55 vs. 7.59). The rotation is genuine, and the mechanism is geometrically principled. In magnitude the advantage is small (about 2% of baseline NLL at $C=4$, less at larger C). The extra cost is $d_{\text{model}}^2 + d_{\text{model}}$ parameters and no observable wall-clock overhead, so when the residual-stream width is large enough that this is a small fraction of the model it is worth including by default; in tighter parameter budgets, skipping it costs $\sim 2\%$ NLL—a real but modest amount.

The C -dependent collapse of the gap (from $\Delta = 0.054$ at $C=4$ to $\Delta = 0.008$ at $C=16$) then admits a natural reading: W_{pos} must rotate P out of a C -dimensional channel span using a fixed d_{model}^2 parameter budget, and the rotation becomes increasingly over-constrained as C grows. We cannot fully separate this from other factors at large C (driver/distractor pressure shifting representational effort elsewhere; larger seed-level optimisation variance), but it is the mechanism-internal explanation.

Channel-independent (ci). CI runs the transformer once per channel and combines the per-channel hidden streams in a single head. It underperforms throughout. Two things contribute. (i) The head is the structural bottleneck: C separate streams have to be combined linearly into a single prediction, which loses the cross-channel interaction that the synthetic task is built on. (ii) It *overfits universally* at the main-sweep scale—all 20 seeds at every $C \in \{4, 8, 16\}$ flag “overfitting”, best val NLL is reached at epoch ~ 40 – 60 , and val NLL drifts upward by 0.12–0.29 nats afterwards. The CI NLL we report is the effective-early-stopping minimum; without early stopping the encoder would look substantially worse. The wall-clock cost is $\sim 8\times$ LINEAR at $C=16$. On this evidence CI is dominated by the per-channel- W_k alternatives.

Channel-as-token (cat). CAT is the more interesting case. It is well converged (all 20 seeds at $C \in \{4, 8\}$ flag “converged”, best epoch ~ 280 – 300 , final-equals-best to four decimals), so the synthetic-benchmark gap to the per-channel- W_k tier (2.35 vs. ~ 2.15 NLL at $C=4$) is a property of its converged optimum, not a convergence artefact. On ETTh1, where all seven variates are informative measurements of one physical system, that gap closes and CAT posts the lowest mean NLL. Paired-difference tests at 20 seeds (Section 3.12) place CAT decisively above MLP ($p=1.4\times 10^{-7}$) and LINEAR-PPE ($p=3\times 10^{-4}$), marginally above CONCAT ($p=0.002$), and *statistically tied* with LINEAR ($p=0.14$) and LINEAR-ORTHO ($p=0.10$); best-bin accuracies in the top tier are indistinguishable. The practical reading is that CAT is competitive on ETTh1, not clearly best. The architecture cost is real on top of that: CAT multiplies the sequence length by C and the attention cost by C^2 , so even at our small C it runs $\sim 5\times$ LINEAR at $C=4$ and $\sim 17\times$ at $C=8$, and the operating context length the GPU has to hold grows in proportion. Reaching for CAT is justified only when the task is known to be one where every channel is a genuine informative measurement of one underlying system and the cross-variate attention can pay—an a priori judgment about the data, not something this paper resolves.

No convergence-time penalty for involved encoders. MLP, LINEAR-ORTHO, and LINEAR-PPE reach within 0.05 NLL of their final value at the same trace point as LINEAR (~ 150 epochs) and cost the same per-epoch wall time. Within the family that uses per-channel W_k on a shared backbone, encoder choice does not change either the convergence schedule or compute budget.

Reading for non-numerical inputs. The mechanism is geometric and encoding-agnostic, which suggests one natural extension worth flagging: when the per-channel inputs are precomputed embeddings from an outside source (e.g. text, image, or tabular feature embeddings fed into a multi-stream transformer), the LINEAR-PPE-style learned rotation has a ready analog as a per-stream rotation that pushes precomputed embedding subspaces off the positional subspace and off one another. Whether this gives gains comparable to the $\sim 2\%$ NLL improvement we measure here, or larger because external embeddings carry more structured competing geometry, is an interesting question for another study.

Practical near-equivalence within the top tier. The 20-seed paired tests reported above are powerful enough to resolve several small inter-encoder gaps, and it is worth separating that statistical power from a practical claim. The absolute differences in NLL among the per-channel- W_k family (LINEAR, LINEAR-ORTHO, CONCAT, MLP, LINEAR-PPE) at $C=4$ span 0.02 NLL on a baseline of 2.16; LINEAR-PPE’s lead over LINEAR is 0.041 NLL, about 2% of the baseline; MLP’s lead over the linear family at $C=16$ is 0.011–0.025 in the data-rich regime. These are real gaps under our experimental design, but they are small relative to the spread induced by typical task-design choices (channel count, sequence length, label binning, etc.). We would not expect them to translate cleanly to practically significant differences on new tasks with new data distributions. Within the top tier the encoder choice is, in any practical sense, close to a free parameter, and for the sake of simplicity we recommend defaulting to plain LINEAR unless the task at hand gives a compelling reason to do otherwise.

5 Limitations

- The synthetic benchmark deliberately makes channel identity informative. On tasks with exchangeable channels, SUM may not be penalised.
- Main sweep at $d_{\text{model}}=64$, $C \in \{4, 8, 16\}$; ETTh1 at $d_{\text{model}}=56$, $C=7$. The d_{model} sweep covers all top-tier encoders plus SUM at $d \in \{64, 128, 256\}$, $C=4$, but all encoders overfit at $d \geq 128$ under the main-sweep training data (Table 4); the reported best-NLL numbers there are effective-early-stopping minima rather than converged values.
- The main sweep at $N_{\text{series}}=512$ is data-limited. $10\times$ training data improves all top-tier encoders by ~ 0.4 NLL at $C=16$ (Table 3); a similar gap exists at $C=4$ ($2.17 \rightarrow 1.85$, see Section 3.6). The relative ordering of encoders is robust to data scale, but absolute inter-encoder gaps shrink $\sim 5\times$ at $10\times$ data. Readers should treat the $N=512$ inter-encoder gaps as upper bounds on the data-rich-regime gaps.
- CAT is skipped at $C=16$ on the synthetic benchmark for compute reasons ($\mathcal{O}((CT)^2)$ attention).
- LINEAR-PPE was tested at $C \in \{4, 8, 16\}$ on the synthetic benchmark and at $C=7$ on ETTh1.
- Headline comparisons are at 20 paired seeds. A few diagnostics (linear probing, channel masking, channel-bias ablation, MLP first-layer geometry, MSE loss-family check) remain at 5 seeds, labelled as such in their tables.

6 Related work

Channel-as-token and channel-independent TSF. iTransformer (5) treats each variate as a token; PatchTST (4) runs independent backbones per channel; Crossformer (6) does both. These preserve per-channel capacity at the cost of longer sequences or per-channel compute. Our encoders share one token per time step.

Linear input projections in TSF. Many multivariate transformers use a linear stem (2, 3). We isolate this bare projection and compare it explicitly against the shared-scalar alternative and architectural baselines.

Orthogonality regularisation. Bansal et al. (7) use soft-orthogonality penalties for feature decorrelation. We apply the same idea to per-channel input projections and find it unnecessary: the task loss produces the same geometry.

7 Conclusion

On tasks where channel identity is informative, a per-channel linear projection (`nn.Linear(C,  $d_{\text{model}}$ )`) matches every more elaborate encoder we tested up to small, statistically real but practically modest, differences. The task loss drives the channel projections to near-orthogonality, and learned norms are larger for outcome-relevant channels. A shared-scalar projection has a capacity-independent ceiling. Nonlinearity, block partitioning, orthogonality penalties, and channel-independent architectures all end up close to the linear baseline: paired-difference analysis at 20 seeds resolves several small inter-encoder gaps as statistically clear, but the gaps span a 0.02–0.04 NLL range on a 2.16 NLL baseline—small enough that we would not expect them to translate cleanly to practically significant differences on new tasks with new data distributions.

Three variants deserve a footnote rather than a full carve-out, with caveats:

- LINEAR-PPE consistently leads the linear family at every C we tested (paired p from 2.6×10^{-6} to 0.036, gap shrinking from 0.041 to 0.016 NLL). A direct geometric probe shows the mechanism is positional-channel orthogonalisation rather than positional subspace compression. The rotation has a clean geometric justification and adds $d_{\text{model}}^2 + d_{\text{model}}$ parameters with no observable wall-clock cost.
- MLP pulls ahead of the linear family at $C=16$ under paired analysis (paired $p \in [3 \times 10^{-4}, 5 \times 10^{-12}]$ vs. LINEAR, LINEAR-ORTHO, CONCAT; gap shrinks $\sim 3\times$ at $10\times$ training data but stays statistically clear). It is worse than LINEAR at $C=4$, and worse than the linear family at $C=4$, $d_{\text{model}} \geq 128$ where it overfits hardest.
- CAT posts the lowest mean NLL on ETTh1, separating cleanly from MLP and LINEAR-PPE under paired tests but *statistically tied* with LINEAR and LINEAR-ORTHO ($p > 0.1$). It costs $\sim 5\times$ wall clock at $C=4$ and substantially more at higher C .

Practical recommendation. For simplicity, default to `nn.Linear(C,  $d_{\text{model}}$ )`. Within the per-channel- W_k family the encoder choice is, for any practical purpose, close to a free parameter, and adopting the simplest choice unless something specific in the task tips the balance is the path of least surprise. If the residual-stream width makes d_{model}^2 extra parameters cheap, projecting the sinusoidal positional encoding through a learned linear layer (LINEAR-PPE) is a justified extra: it

has a clean geometric mechanism, costs no wall-clock overhead, and gains a small but consistent $\sim 2\%$ NLL at $C=4$. Reach for something heavier (MLP, CAT) only when there is a real and strong reason in the task at hand—e.g. a high-channel regime where the MLP nonlinearity may genuinely gate informative channels against distractors, or a setting where CAT’s $\sim 5\times$ wall-clock cost and $C\cdot T$ context length are justified by a known cross-variate structure.

References

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *NeurIPS*, 2017.
- [2] H. Zhou, S. Zhang, J. Peng, S. Zhang, J. Li, H. Xiong, and W. Zhang. Informer: beyond efficient transformer for long sequence time-series forecasting. In *AAAI*, 2021.
- [3] H. Wu, T. Hu, Y. Liu, H. Zhou, J. Wang, and M. Long. TimesNet: temporal 2D-variation modeling for general time series analysis. In *ICLR*, 2023.
- [4] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam. A time series is worth 64 words: long-term forecasting with transformers. In *ICLR*, 2023.
- [5] Y. Liu, T. Hu, H. Zhang, H. Wu, S. Wang, L. Ma, and M. Long. iTransformer: inverted transformers are effective for time series forecasting. In *ICLR*, 2024.
- [6] Y. Zhang and J. Yan. Crossformer: transformer utilizing cross-dimension dependency for multivariate time series forecasting. In *ICLR*, 2023.
- [7] N. Bansal, X. Chen, and Z. Wang. Can we gain more from orthogonality regularizations in training deep networks? In *NeurIPS*, 2018.